

Approximate Aggregation

Outline

- Introduction to Approximate Query Processing(AQP)
- Related Work on AQP (over Joins)
- What about Non-Foreign-Key Joins?
- Weighted Sampling for AQP



Introduction to AQP

Approximate Query Processing

- Primarily for aggregate queries
 - count, sum, avg
- Exact answers are **NOT** *always* required
 - Exploratory queries in Decision Support Systems(DSS)
 - Interactive response times (in several seconds)
 - Big data visualization
 - Incomplete or noisy data
- Answering queries based on **samples**
 - Sampling can handle more general queries
 - Although other techniques exist: histograms, wavelets

```
SELECT sum(price)
FROM orders
WHERE region = 'Asia'
```

Exact: \$12345

*AQP : \$12300 ± 100
with 95% confidence*

Sampling for AQP

- A relatively small random sample can well-represent the entire data
- Example: approximate answer for *SUM()* queries
 - e.g., S is a 1% sample of *orders*
 - Q : *SELECT sum(price) FROM orders WHERE region = 'Asia'*
 - X = the result of evaluating Q on S
 - Final answer = $(X * 100) \pm \Delta$ with $xx\%$ confidence
- Confidence Intervals
 - Most widely used: CLT-based confidence interval
 - Central Limit Theory
 - $\{X_1, X_2, \dots, X_n\}$: a sequence of i.i.d. random variables
 - $E[X_i] = \mu, Var[X_i] = \sigma^2, S_n = (X_1 + X_2 + \dots + X_n) / n$
 - $\sqrt{n} * (S_n - \mu)$ converge to $N(0, \sigma^2)$

Related Work

Notation & Terminology

- Grouping/Aggregate Columns
 - e.g., *Q*: `SELECT sum(price) FROM orders GROUP BY region`
 - grouping column: *region*, aggregate column: *price*



Sampling Methods

- Stratified Sampling
 - divide the population into homogeneous subgroups
 - simple random sampling in each subgroup
- Cluster Sampling
 - divide the population into clusters
 - random sampling the clusters(, then the elements in them)
- PPS Sampling
 - probability-proportional-to-size sampling

$$\hat{\tau} = \frac{1}{n} \sum_{i=1}^n \frac{y_i}{p_i}, \quad \text{var}(\hat{\tau}) = \frac{1}{n} \sum_{i=1}^N p_i \left(\frac{y_i}{p_i} - \tau \right)^2$$



Cardinality Estimation

- CE is a traditional research area of DBMS
- Equivalent to *COUNT*(*) queries in AQP
- Similar techniques
 - histogram, wavelet, sketch, sampling
 - sampling is not widely adopted due to its significant overhead
 - recent research shows the benefits of sampling in join size estimation



Uniform Sampling over Joins

- For hard cases, uniform sampling on either side is infeasible
- Attach weights based on the frequency of join attribute values
 - weights proportional to the frequency in the other side
- Consider $T_1 \bowtie T_2$
 - Let $F_i(a_j)$ be the frequency of join attribute value a_j in table T_i
 - S_1 : weighted sampling on T_1
 - pick each (a_1) with probability proportional to $F_2(a_1) = 1$
 - pick (a_2) with probability proportional to $F_2(a_2) = n$
 - For each tuple in S_1 :
 - choose one joinable tuple from T_2 randomly

T_1		T_2	
A		A	B
a_1		a_1	1
a_1		a_2	2
...		a_2	3
a_1	
a_2		a_2	$n+1$

AQP over a single relation

- **Value skew** in aggregate columns
 - precompute $SUM(column)$ and **scale factors** of tuples
 - outlier index
- **Frequency skew** among groups
 - congressional sampling: divide the sample size equally among groups
 - small group sampling: take all tuples of small groups into the sample
 - stratified sampling on grouping columns
- Highly selective filter conditions
 - workload-based weighted/stratified sampling
 - auxiliary index
- Non-correlated error among groups
 - precompute $SUM(column)$



AQP over Joins

- Extend methods on a single relation to foreign-key joins
 - draw a sample from the fact table, then join it with the dimension table
- Draw simple random samples from both sides
 - ripple join
- Index-based correlated sampling
 - select random tuples of one side
 - sample tuples of the other side that can join with them
- Cluster sampling over join attributes
 - take simple random sample of join attribute values (using hash)
 - keep tuples whose join attribute values have been sampled



Non-Foreign-Key Joins?

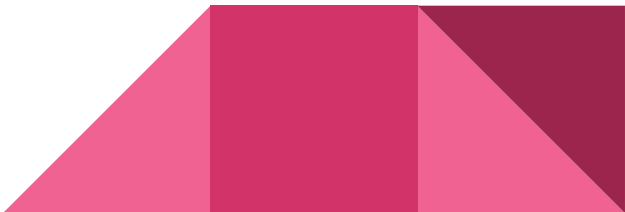
Many-to-Many Joins

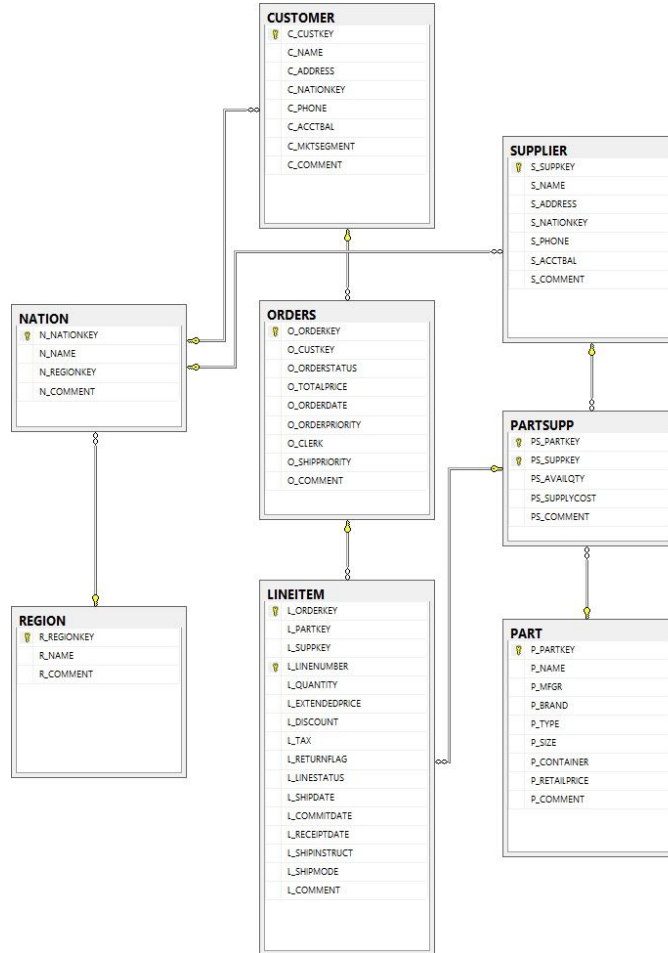
- Most joins are foreign-key joins (star/snowflake schema)

- Self-join

```
SELECT    avg(abs(A.salary - B.salary))  
FROM      employee A  
JOIN      employee B USING (department)  
WHERE     A.id < B.id  
GROUP BY department
```

- What are the **semantics** of aggregations after many-to-many joins?





Weighted Sampling for AQP

Value-Weighted Sampling

- Idea: *SUM()* can be converted to *COUNT(*)*
 - e.g., *SUM(b)*
 - convert table *T* into a virtual table *T_v*
 - for each tuple *t*, repeat itself *t.b* times and set *t.b = 1*

T		T _v	
a	b	a	b
a1	1	a1	1
a2	2	a2	1
a3	3	a2	1
		a3	1
		a3	1
		a3	1

Value-Weighted Sampling

- Idea: *SUM()* can be converted to *COUNT(*)*

- e.g., *SUM(b)*
- convert table *T* into a virtual table *T_v*
- for each tuple *t*: repeat *t.b* times & set *t.b = 1*

T			T _v		
p	a	b	a	b	p
1/6	a1	1	a1	1	1/6
2/6	a2	2	a2	1	1/6
3/6	a3	3	a2	1	1/6
			a3	1	1/6
			a3	1	1/6
			a3	1	1/6

- Weighted sampling in *T*
 - equivalent to **uniform sampling** in *T_v*
 - attach each tuple *t* a weight proportional to *t.b*

Drawbacks

- Aggregate Expressions

- `sum(l_extendedprice * (1 - l_discount) * (1 + l_tax))`
- column-stores may materialize these expressions

- Multi-aggregation

```
SELECT
    sum(l_quantity) as sum_qty,
    sum(l_extendedprice) as sum_base_price,
    sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
    sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
    avg(l_quantity) as avg_qty,
    avg(l_extendedprice) as avg_price,
    avg(l_discount) as avg_disc
FROM lineitem ...
```

- Negative values

Frequency-Weighted Sampling

- Consider $T_1 \bowtie T_2$
 - Let $F_i(a_j)$ be the frequency of join attribute value a_j in table T_i
 - S_1 : weighted sampling on T_1
 - pick each (a_1) with probability proportional to $F_2(a_1) = 1$
 - pick (a_2) with probability proportional to $F_2(a_2) = n$
 - For each tuple in S_1 :
 - choose one joinable tuple from T_2 randomly


T_1		T_2	
A		A	B
a_1		a_1	1
a_1		a_2	2
...		a_2	3
a_1	
a_2		a_2	$n+1$

Value-Weighted & Frequency-Weighted

- Star/snowflake Schema
- Value-weighted: draw samples from the fact table
 - **aggregate columns** belong to the fact table
- Frequency-weighted: draw samples from dimension tables
 - fact \rightarrow dimension: **many-to-one**
- Frequency-weighted + value-weighted?
 - t_d : dimension tuple, f : fact table, j : join attribute, a : aggregate attribute

$$weight(t_d) = \sum_{i=0}^{F_f(t_d, j)} t_{f, a}^i$$

$$weight(t_f^i) = \frac{t_{f, a}^i}{weight(t_d)}$$



d.j	f.j	f.a
y_1	y_1	1
y_1	y_1	2
y_1	y_1	3
y_1	y_1	4

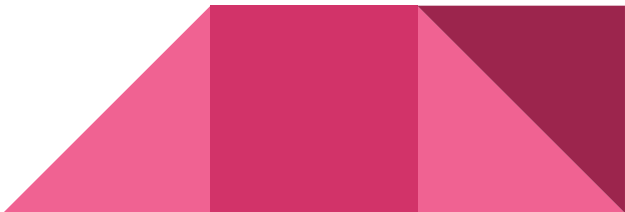
$$weight(t_d) = \sum_{i=0}^{F_f(t_d.j)} t_{f.a}^i$$

$$weight(t_f^i) = \frac{t_{f.a}^i}{weight(t_d)}$$

Cost Model for AQP?

- Uniform samples or weighted samples?
- Draw samples from fact tables or from dimension tables?
- Sampling-selection or selection-sampling?

```
SELECT sum(orders.price)
FROM   orders, customer
WHERE  orders.cid = customer.id
      AND customer.name = 'Bob'
```

- Sampling as an operator
 - Factors: sample size, selectivity
- 

Backup Slides

Workload-Weighted Sampling

- Intuition: many queries only touch a portion of the data
 - latest events, constant predicts, ...

```
SELECT count(*)  
FROM   events AS ev  
WHERE  ev.timestamp - current_timestamp  
       < INTERVAL '24 hours '
```

- Idea: maintain a reservoir sample for **latest seen** tuples
- Idea: attach weights based on frequency of **access** to each tuple



Workload-Based AQP

- Analyze the given workload
 - Maintain different types of samples
 - uniform sample, time-weighted sample, workload-weighted sample, ...
 - Classify incoming queries to use these samples
-
- Public Query Logs: Sloan Digital Sky Survey (www.sdss.org)
 - 4,000,000+ query logs per month
 - very few aggregate queries (0.01%) 😞