





Graph Neural Networks

杨帆

2018年7月10日

Motivations for This Talk (1)

- (My) ultimate goal: to build models that accurately **predict the future**
 - We have already learned some powerful sequential models
 - E.g., RNNs, stochastic RNNs, WaveNet, state space models
 - They are good at modeling a set of **independent** sequences
 - Real-world sequences are often **interacting** with each other
 - Leveraging such interactions can improve the predictive ability of models



Motivations for This Talk (2)

- **Simulators** are perfect tools for forecasting dynamical systems
 - However, building simulators for real environment is not easy
 - Requiring a lot of domain-specific knowledge, e.g., physics
- There is a latest trend toward **learning** simulators from observations
 - Neural Physics Engines, World Models
- Graph NNs are the key building block behind such learned simulators
 - Do well in modeling the relations between objects, while remaining **tractable**
- Besides, Graph NNs are also promising tools for:
 - Representation learning on graphs, i.e., node/graph embedding
 - Generalizing deep models to graph-structured data
 - Semi-supervised learning
 - Model-based reinforcement learning

Outline

- Introduction to Graph Neural Networks
- Hierarchical Pooling in GNN
- Modeling Interacting Systems with GNN
- Conclusions and Further Discussion

Outline

- **Introduction to Graph Neural Networks**
- Hierarchical Pooling in GNN
- Modeling Interacting Systems with GNN
- Conclusions and Further Discussion

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. M Defferrard, et al. NIPS 2016.

Semi-Supervised Classification with Graph Convolutional Networks. TN Kipf, et al. ICLR 2017.

Geometric Deep Learning: Going beyond Euclidean data. MM Bronstein, et al. IEEE Signal Processing Magazine, 2017.

Graph Signal Processing: Overview, Challenges, and Applications. A Ortega, et al. Proceedings of IEEE, 2018.



Representation Learning on Graphs: Methods and Applications. WL Hamilton, et al. Arxiv 2018.

Relational inductive biases, deep learning, and graph networks. PW Battaglia, et al. Arxiv 2018.

The Need for Graph Neural Networks

- Graph-structured data are ubiquitous


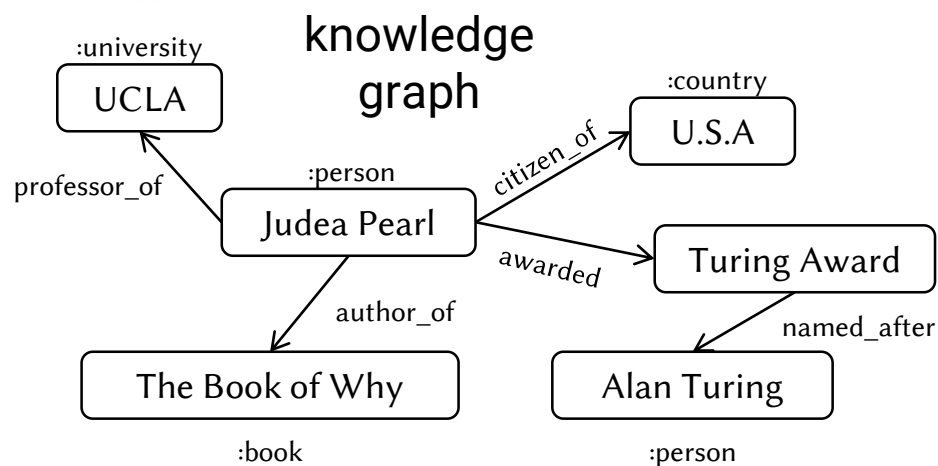
social network
citation network
similarity graph



brain network




roadmaps





molecules

The Need for Graph Neural Networks


- Consider the basic graph classification task



- For image classification, we use **convolutions** to extract features

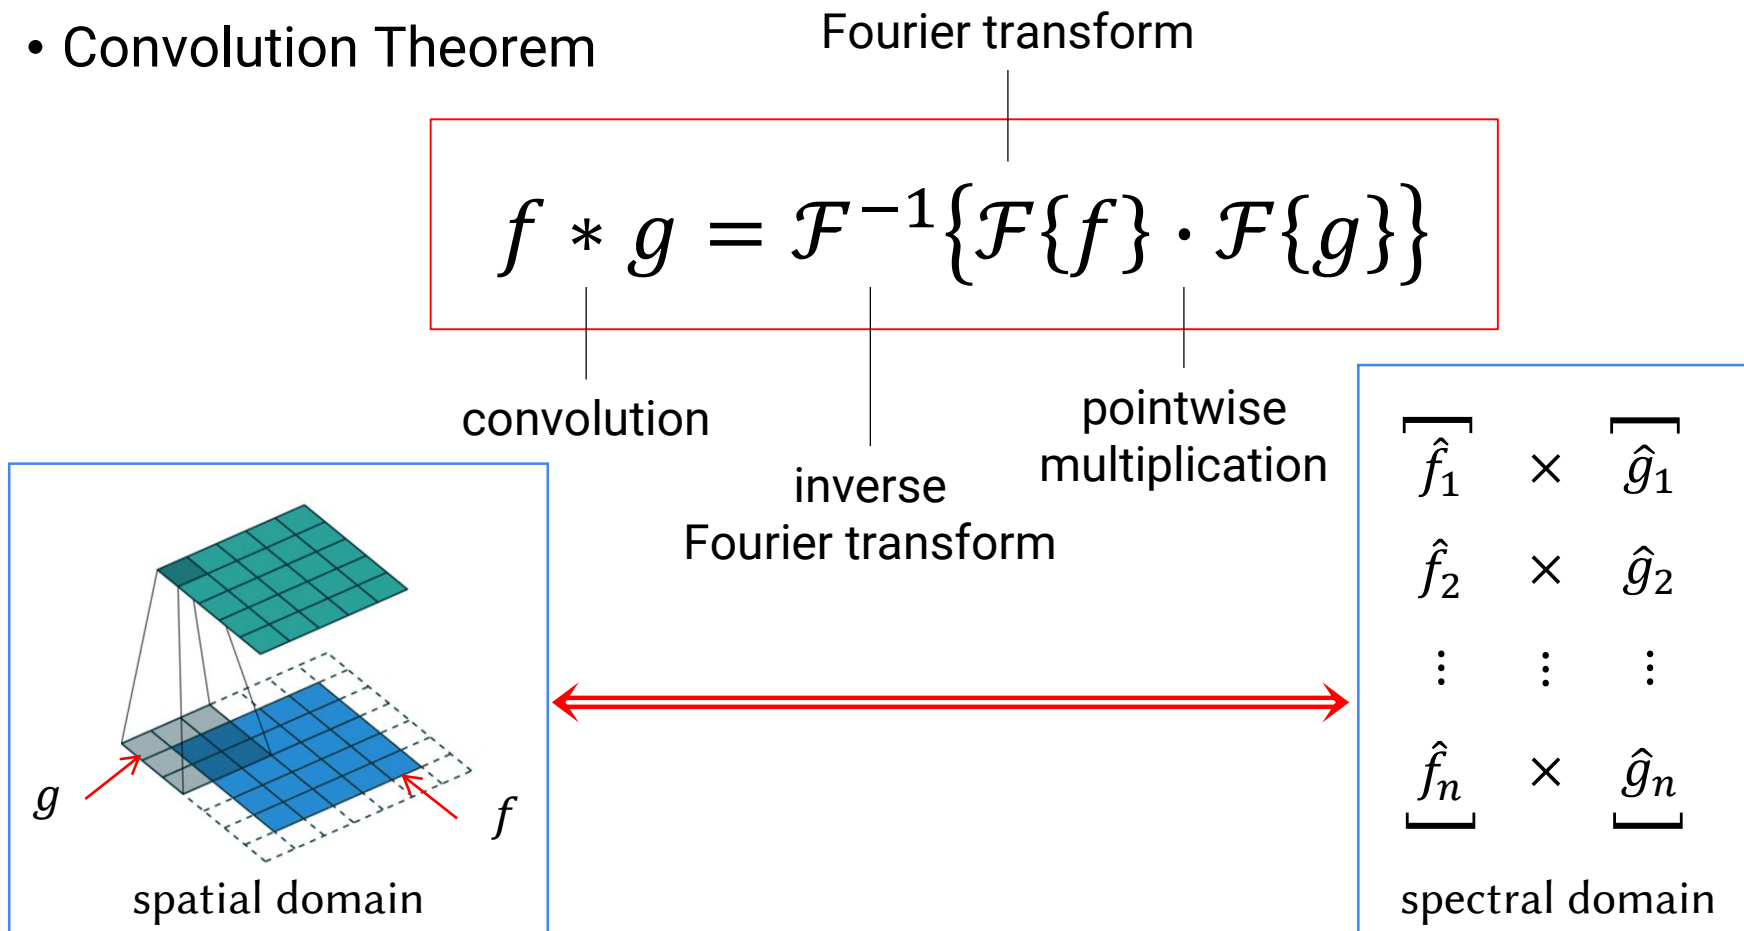


How to extract features from **irregular** graphs?



The Two Faces of Convolutions

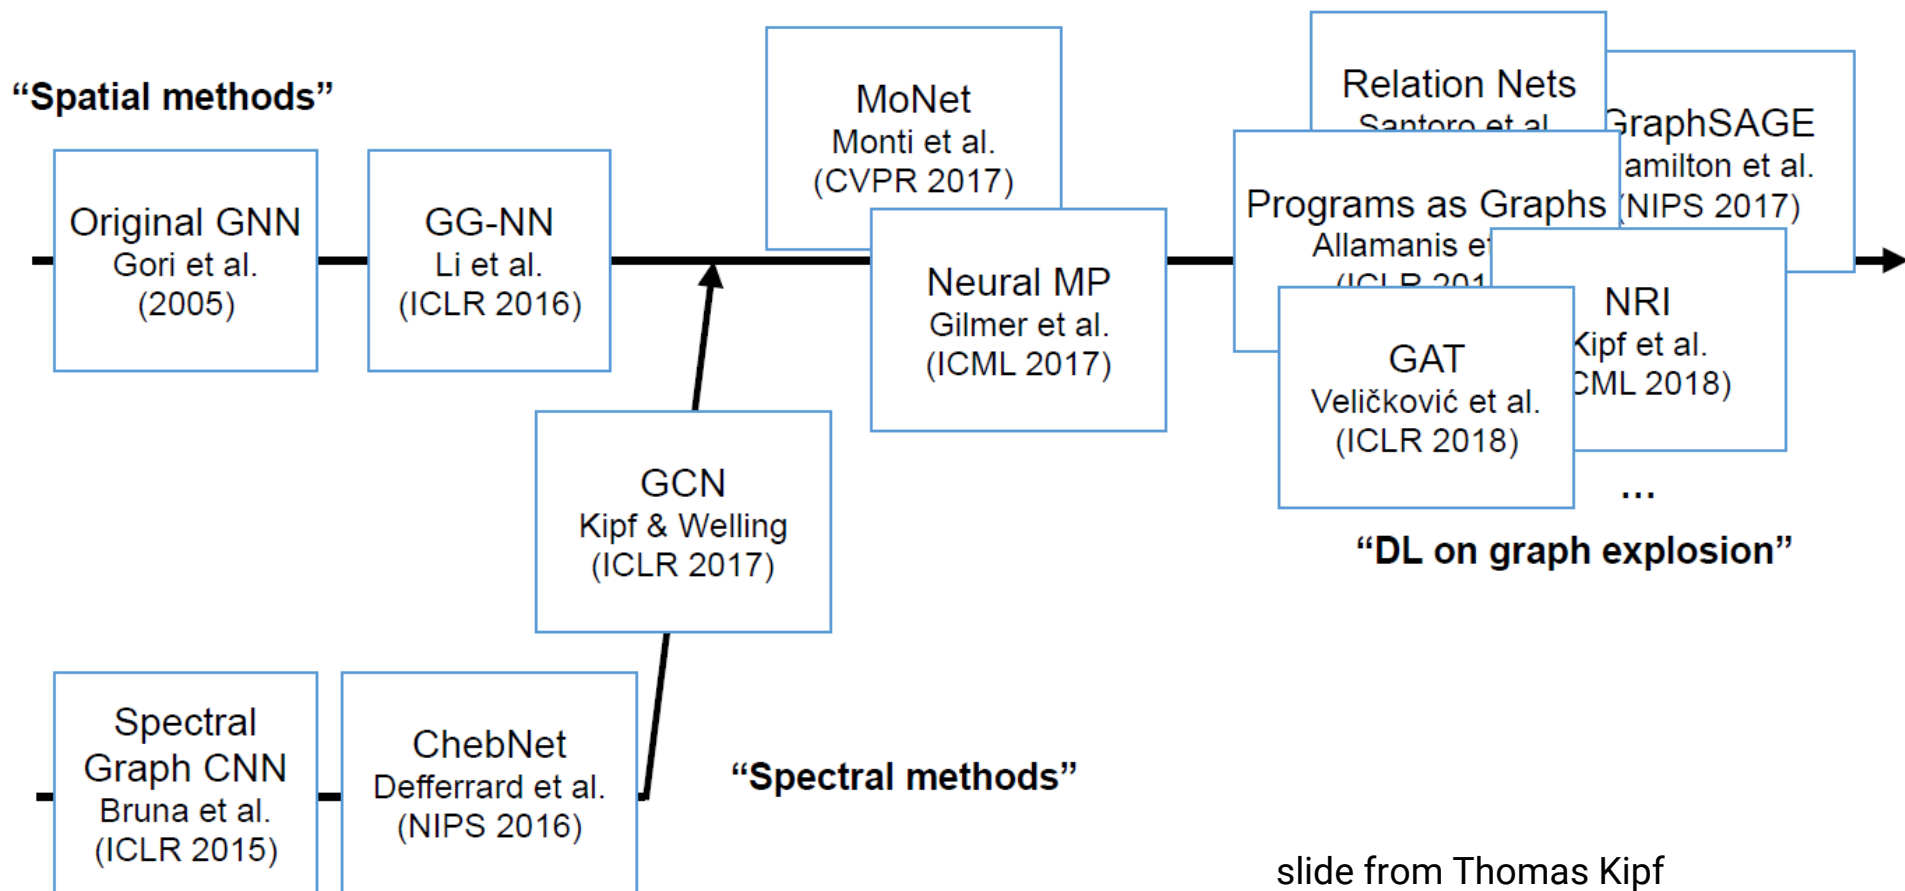
- Convolution Theorem



- Convoluting with a filter \Leftrightarrow cross-correlating with the reverse filter
 - i.e., *sliding inner-product*

History of Graph Neural Networks

- GNNs were also developed from two perspectives



slide from Thomas Kipf
<http://tkipf.github.io/misc/SlidesCambridge.pdf>

Spectral Graph Theory Basics

- Spectral graph convolution is built on top of **spectral graph theory**
 - Study the properties of a graph through the lens of its **Laplacian matrix**
 - Or adjacency matrix

- $G = (V, E)$ is a graph
 - undirected, unweighted, no self-loops
 - $N = \# \text{vertices}$

- **adjacency matrix** A is $N \times N$ matrix

$$A_{ij} = \begin{cases} 1 & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}$$

- **degree matrix** D is $N \times N$ matrix

$$D_{ij} = \begin{cases} d_i & \text{if } i = j \\ 0 & \text{otherwise} \end{cases}$$

- $d_i = \sum_j A_{ij}$ is the degree of vertex i


- **Laplacian matrix** $L = D - A$

A

	1	2	3	4	5
1	0	1	0	1	0
2	1	0	1	0	0
3	0	1	0	1	1
4	1	0	1	0	0
5	0	0	1	0	0

D

	1	2	3	4	5
1	2	0	0	0	0
2	0	2	0	0	0
3	0	0	3	0	0
4	0	0	0	2	0
5	0	0	0	0	1



L

	1	2	3	4	5
1	2	-1	0	-1	0
2	-1	2	-1	0	0
3	0	-1	3	-1	-1
4	-1	0	-1	2	0
5	0	0	-1	0	1

Spectral Graph Theory Basics

- The weighted and normalized definition

$$A_{ij} = \begin{cases} w_{ij} & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}, \quad D_{ij} = \begin{cases} \sum_j A_{ij} & \text{if } i = j \\ 0 & \text{otherwise} \end{cases}$$

$$L = D^{-1/2}(D - A)D^{-1/2} = I - D^{-1/2}AD^{-1/2}$$

- Laplacian matrix is positive-semidefinite and can be diagonalized as:

$$L = U\Lambda U^T = \sum_{i=1}^N \lambda_i \mathbf{u}_i \mathbf{u}_i^T$$

where

$$U = [\mathbf{u}_1 \quad \mathbf{u}_2 \quad \dots \quad \mathbf{u}_N] \in \mathbb{R}^{N \times N}$$

$$\Lambda = \text{diag}([\lambda_1, \lambda_2, \dots, \lambda_N]) \in \mathbb{R}^{N \times N}$$

Graph Signal Processing

- 1D signal: $x \in \mathbb{R}^N$
 - x_i is the value at time i

- Fourier transform

$$\hat{x} = Fx$$

- spectral components

$$\left\{ F_k[n] = e^{-j\frac{2\pi}{N}kn} / \sqrt{N} : n = 0, \dots, N-1 \right\}_{k=0}^{N-1}$$

- frequencies

$$2\pi k/N, k = 0, 1, \dots, N-1$$

- convolution

$$h * x = F^{-1}((Fh) \cdot (Fx))$$

- 1D graph signal: $x \in \mathbb{R}^N$
 - x_i is the value at node i

- graph Fourier transform

$$\hat{x} = U^T x$$

- spectral components

$$\{\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_N\}$$

- graph frequencies

$$\lambda_1, \lambda_2, \dots, \lambda_N$$


- graph convolution

$$h_\theta * x \triangleq U(\theta \cdot (U^T x))$$

diagonalization: $O(N^3)$, matmul: $O(N^2)$ 😞

Spatial Graph NNs


- Spatial Graph NNs are based on the idea of **message passing**
 - Some GNNs were inspired by the spatial interpretation of convolutions



$$h_{p_i}^{l+1} = \sum_{p_j \in \mathcal{N}(p_i)} W_j^l h_{p_j}^l$$

where

$$\mathcal{N}(p) = \text{neighbors}(p)$$



- Simple message passing on general undirected graphs

Input: Graph $G = (V, E)$, non-linearity σ , differentiable AGGREGATE and COMBINE, weights \mathbf{W} , node features $\{h_v^l: v \in V\}$


Output: new node features $\{h_v^{l+1}: v \in V\}$

for $v \in V$ **do**

$h_{\mathcal{N}(v)}^{l+1} \leftarrow \text{AGGREGATE}(\{h_u^l: u \in \mathcal{N}(v)\})$


$h_v^{l+1} \leftarrow \sigma(\mathbf{W} \cdot \text{COMBINE}(h_v^l, h_{\mathcal{N}(v)}^{l+1}))$

end



The Graph Network Framework

- **Graph Network** [Battaglia 2018] generalizes and extends various GNNs
 - They use a generalized definition of graph $G = (\mathbf{u}, V, E)$




- Directed: one-way edges $\mathbf{e}_k = (\mathbf{v}_{s_k}, \mathbf{v}_{r_k})$
- Attributed: vertices and edges have attributes
- Attribute: encoded as a vector or set
- Global attribute: a graph-level attribute \mathbf{u}
- Multi-graph: there can be more than one edge between vertices, including self-edges

- A GN block takes as input a graph $G = (\mathbf{u}, V, E)$


- And outputs a new graph $G' = (\mathbf{u}', V', E')$

- Differentiable and composable




Insides of a Graph Network Block


- A GN block: 3 update functions ϕ and 3 aggregation functions ρ
 - Edge update \rightarrow node update \rightarrow global update




(a) Edge update



(b) Node update



(c) Global update



$$\mathbf{e}'_k = \phi^e(\mathbf{e}_k, \mathbf{v}_{s_k}, \mathbf{v}_{r_k}, \mathbf{u})$$

$$\bar{\mathbf{e}}'_i = \rho^{e \rightarrow v}(E'_i)$$

$$\mathbf{v}'_i = \phi^v(\bar{\mathbf{e}}'_i, \mathbf{v}_i, \mathbf{u})$$

$$\bar{\mathbf{e}}' = \rho^{e \rightarrow u}(E')$$

$$\bar{\mathbf{v}}' = \rho^{v \rightarrow u}(V')$$

$$\mathbf{u}' = \phi^u(\bar{\mathbf{e}}', \bar{\mathbf{v}}', \mathbf{u})$$

Configuring Within-Block Structure

- Note that a GN block is a computation flow abstraction
- Functions ϕ and ρ must be instantiated with concrete implementation
- Example:

$$\phi^e(\mathbf{e}_k, \mathbf{v}_{s_k}, \mathbf{v}_{r_k}, \mathbf{u}) := \text{NN}_e([\mathbf{e}_k, \mathbf{v}_{s_k}, \mathbf{v}_{r_k}, \mathbf{u}])$$

$$\phi^v(\bar{\mathbf{e}}'_i, \mathbf{v}_i, \mathbf{u}) := \text{NN}_v([\bar{\mathbf{e}}'_i, \mathbf{v}_i, \mathbf{u}])$$

$$\phi^u(\bar{\mathbf{e}}', \bar{\mathbf{v}}', \mathbf{u}) := \text{NN}_u([\bar{\mathbf{e}}', \bar{\mathbf{v}}', \mathbf{u}])$$

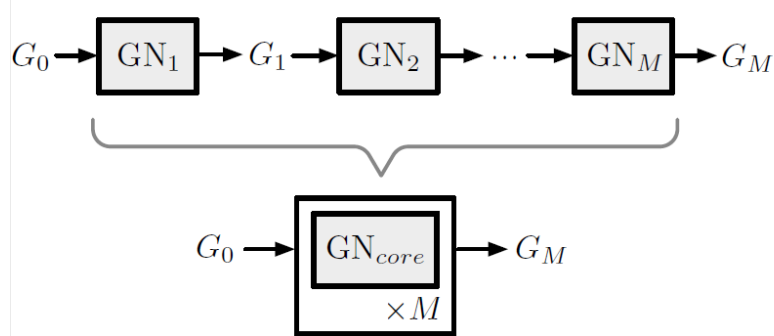
$$\rho^{e \rightarrow v}(E'_i) := \sum_k \mathbb{I}(r_k = i) \mathbf{e}'_k$$

$$\rho^{v \rightarrow u}(V') := \sum_i \mathbf{v}'_i, \quad \rho^{e \rightarrow u}(E') := \sum_k \mathbf{e}'_k$$


- Other aggregation functions can be used, e.g., average/max

Compositions of Graph Network Blocks


- GN blocks can be stacked together to build more power models




(a) Composition of GN blocks



(b) Encode-process-decode



(c) Recurrent GN architecture



Outline

- Introduction to Graph Neural Networks
- **Hierarchical Pooling in GNN**
- Modeling Interacting Systems with GNN
- Conclusions and Further Discussion

Neural Expectation Maximization. K Greff, et al. NIPS 2017


Relational Neural Expectation Maximization: Unsupervised Discovery of Objects and their Interactions. S Steenkiste, et al. ICLR 2018

Hierarchical Graph Representation Learning with Differentiable Pooling. R Ying, et al. Arxiv 2018



Flexible Neural Representation for Physics Prediction. D Mrowca, et al. Arxiv 2018

The Need for Hierarchical Pooling/Clustering

- Graph Networks are flat: the input and output have the same #vertices
- We would like to introduce some hierarchical structure to:
 - Aggregate node features in a hierarchical way




- Model the interactions at a coarser level
- Make the message passing more efficient




Hard Pooling

- Sometimes the hierarchical structure is fixed or relatively stable
- We can use hierarchical (graph) clustering algorithms to build a tree



- Hierarchical Graph Convolution





message passing:

- (1) ϕ^{L2A} : leave $L \rightarrow$ ancestors A
- (2) ϕ^{WS} : between siblings
- (3) ϕ^{A2D} : $A \rightarrow$ descendants D


Soft Pooling

- In many cases, the hierarchical structure is not fixed
 - Evolving over time
 - Different hierarchies across graphs



- Possible solution: learn the **assignment matrices**
 - Interestingly, an assignment matrix itself can be computed using a GNN

$$\begin{aligned} Z^{(l)} &= \text{GNN}_{l,\text{embed}}(A^{(l)}, X^{(l)}) \in \mathbb{R}^{n_l \times d} \\ S^{(l)} &= \text{softmax}\left(\text{GNN}_{l,\text{pool}}(A^{(l)}, X^{(l)})\right) \in \mathbb{R}^{n_l \times n_{l+1}} \\ X^{(l+1)} &= S^{(l)T} Z^{(l)} \in \mathbb{R}^{n_{l+1} \times d} \\ A^{(l+1)} &= S^{(l)T} A^{(l)} S^{(l)} \in \mathbb{R}^{n_{l+1} \times n_{l+1}} \end{aligned}$$






Outline

- Introduction to Graph Neural Networks
- Hierarchical Pooling in GNN
- **Modeling Interacting Systems with GNN**
- Conclusions and Further Discussion

Relational Neural Expectation Maximization: Unsupervised Discovery of Objects and their Interactions. S Steenkiste, et al. ICLR 2018
Neural Relational Inference for Interacting Systems. T Kipf, M Welling. ICML 2018
Graph Networks as Learnable Physics Engines for Inference and Control. A Sanchez-Gonzalez, et al. ICML 2018

Interacting Dynamical Systems

- Most real-world dynamical systems can be decomposed into smaller dynamics that interact with each other



- Given sequential observations of individual objects, we would like to:
 - Infer the latent interaction structure
 - Learn the dynamical model of the interacting system
- So we could understand what happens and better predict the future


Modeling Interacting Systems with GNN

- Assume there are K objects and their corresponding sequential observations $\{x_k^{1:T}\}_{k=1}^K$

$$\tilde{h}_{1:K}^t = \text{GNN}(h_{1:K}^t)$$

$$h_k^{t+1} = \text{RNN}(x_{1:K}^t, \tilde{h}_k^t)$$

$$x_{1:K}^{t+1} | x_{1:K}^t \sim p(\cdot | f(h_{1:K}^{t+1}))$$



where $\tilde{h}_{1:K}^t = \text{GNN}(h_{1:K}^t)$ is implemented as follows:

$$\begin{aligned} \tilde{h}_k^t &= [\hat{h}_k^t, E_k^t], & \hat{h}_k^t &= \text{MLP}^{enc}(h_k^t), & E_k^t &= \sum_{i \neq k} \alpha_{k,i}^t \cdot e_{k,i}^t \\ \alpha_{k,i}^t &= \text{MLP}^{att}(\xi_{k,i}^t), & e_{k,i}^t &= \text{MLP}^{eff}(\xi_{k,i}^t), & \xi_{k,i}^t &= \text{MLP}^{emb}([\hat{h}_k^t, \hat{h}_i^t]) \end{aligned}$$

Building Latent Variable Models with GNN

- **NRI** introduces discrete random variables z_{ij} to represent edge types

$$\mathbf{z} \sim p(\mathbf{z})$$

VAE is used for model learning and inference

$$h_{1:K}^{t+1} = \text{GNN}(h_{1:K}^t, x_{1:K}^t, \mathbf{z})$$

$$\mu_k^{t+1} = x_k^t + f_{\text{out}}(h_k^{t+1})$$


$$x_{1:K}^{t+1} | x_{1:K}^t, \mathbf{z} \sim \mathcal{N}(\mu_{1:K}^{t+1}, \sigma^2 \mathbf{I})$$

where $h_{1:K}^{t+1} = \text{GNN}(h_{1:K}^t, x_{1:K}^t, \mathbf{z})$ is:

$$h_{(i,j)}^t = \sum_{m=1}^M z_{ij,m} f_e^m([h_i^t, h_j^t])$$

$$\text{MSG}_k^t = \sum_{i \neq k} h_{(i,k)}^t$$

$$h_k^{t+1} = \text{GRU}([\text{MSG}_k^t, x_k^t], h_k^t)$$



Neural Relational Inference for Interacting Systems. T Kipf, M Welling. ICML 2018

Outline

- Introduction to Graph Neural Networks
- Hierarchical Pooling in GNN
- Modeling Interacting Systems with GNN
- **Conclusions and Further Discussion**

Conclusions and Further Discussion

- Graph Neural Networks are very likely to be ubiquitous building blocks for deep models
- The theoretical understanding of GNNs is still lacking
 - e.g., spectral graph theory for directed graphs, convergence of message passing, effects of auxiliary edges/channels
- GNNs are useful for modeling interacting dynamical systems
 - There is still room for improvement, e.g., introducing hierarchies, using better models for individual objects

Questions?